

Approximation Algorithms for the Open Shop Problem with Delivery Times

Imed KACEM* and Christophe RAPINE†

Abstract

In this paper we consider the open shop scheduling problem where the jobs have delivery times. The minimization criterion is the maximum lateness of the jobs. This problem is known to be NP-hard, even restricted to only 2 machines. We establish that any list scheduling algorithm has a performance ratio of 2. For a fixed number of machines, we design a polynomial time approximation scheme (PTAS) which represents the best possible result due to the strong NP-hardness of the problem.

Keywords: Scheduling ; Open Shop ; Maximum Lateness ; Approximation ; PTAS

1 Introduction

Problem description. We consider the open shop problem with delivery times. We have a set $\mathcal{J} = \{1, 2, \dots, n\}$ of n jobs to be performed on a set of m machines $\mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3, \dots, \mathcal{M}_m$. Each job j consists of exactly m operations $O_{i,j}$ ($i \in \{1, 2, \dots, m\}$) and has a delivery time q_j , that we assume non negative. For every job j and every index i , operation $O_{i,j}$ should be performed on machine \mathcal{M}_i . The processing time of each operation $O_{i,j}$ is denoted by $p_{i,j}$. At any time, a job can be processed by at most one machine. Moreover, any machine can process only one job at a time. Preemption of operations is not allowed. We denote by $C_{i,j}$ the completion time of operation $O_{i,j}$. For every job j , its completion time C_j is defined as the completion time of its last operation. The lateness L_j of job j is equal to $C_j + q_j$. The objective is to find a feasible schedule that minimizes the maximum lateness L_{\max} , where

*LCOMS, Université de Lorraine, Ile du Saulcy, Metz 57000, France. Contact: imed.kacem@univ-lorraine.fr

†LGIPM, Université de Lorraine, Ile du Saulcy, Metz 57000, France. Contact : christophe.rapine@univ-lorraine.fr

$$L_{\max} = \max_{1 \leq j \leq n} \{L_j\} \quad (1)$$

For any feasible schedule π , we denote the resulting maximum lateness by $L_{\max}(\pi)$. Moreover, L_{\max}^* denotes the maximum lateness of an optimal solution π^* , that is, $L_{\max}^* = L_{\max}(\pi^*)$. According to the tertiary notation, the problem is denoted as $O||L_{\max}$.

Recall that a constant approximation algorithm of performance ratio $\gamma \geq 1$ (or a γ -approximation) is a polynomial time algorithm that provides a schedule with maximum lateness no greater than γL_{\max}^* for every instance. A polynomial time approximation scheme (PTAS) is a family of $(1 + \varepsilon)$ -approximation algorithms of a polynomial time complexity for any fixed $\varepsilon > 0$. If this time complexity is polynomial in $1/\varepsilon$ and in the input size then we have a fully polynomial time approximation scheme (FPTAS).

Related approximation results. According to the best of our knowledge, the design of approximation algorithms has not yet been addressed for problem $O||L_{\max}$. However, some inapproximability results have been established in the literature. For a fixed number of machines, unless $P=NP$, problem $Om||L_{\max}$ cannot admit an FPTAS since it is NP-hard in the strong sense on two machines [8], [9]. The existence of a PTAS for a fixed m is an open question, that we answer positively in this paper. If the number m of machines is part of the inputs, Williamson *et al* [11] proved that no polynomial time approximation algorithm with a performance guarantee lower than $5/4$ can exist, , unless $P=NP$, which precludes the existence of a PTAS. Several interesting results exist for some related problems, mainly to minimize the makespan :

- Lawler *et al* [8]-[9] presented a polynomial algorithm for problem $O2|pmtn|L_{\max}$. In contrast, when preemption is not allowed, they proved that problem $O2||L_{\max}$ is strongly NP-hard, as mentioned above.
- Gonzales and Sahni [4] proved that problem $Om||C_{\max}$ is polynomial for $m = 2$ and becomes NP-hard when $m \geq 3$.
- Sevastianov and Woeginger [10] established the existence of a PTAS for problem $Om||C_{\max}$ when m is fixed.
- Kononov and Sviridenko [7] proposed a PTAS for problem $Oq(Pm)|r_{ij}|C_{\max}$ when q and m are fixed.
- Approximation algorithms have been recently proposed for other variants such as the two-machine routing open shop problem. A sample of them includes Chernykh *et al* [2] and Averbakh *et al* [1].

Finally, we refer to the state-of-the-art paper on scheduling problems under the maximum lateness minimization by Kellerer [6].

Contribution. Unless $P=NP$, problem $Om||L_{\max}$ cannot admit an FPTAS since it is NP-hard in the strong sense on two machines. Hence, the best possible approximation algorithm is a PTAS. In this paper, we prove the existence of such an algorithm for a fixed number of machines, and thus gives a positive answer to this open problem. Moreover, we provide the analysis of some simple constant approximation algorithms when the number of machines is a part of the inputs.

Organization of the paper. Section 2 present some simple preliminary approximation results on list scheduling algorithms. In Section 3, we describe our PTAS and we provide the analysis of such a scheme. Finally, we give some concluding remarks in Section 4.

2 Approximation Ratio of List Scheduling Algorithms

List scheduling algorithms are popular methods in scheduling theory. Recall that a list scheduling algorithm relies on a greedy allocation of the operations to the resources that prevents any machine to be inactive while an operation is available to be performed. If several operations are concurrently available, ties are broken using a priority list. We call a *list schedule* the solution produced by a list scheduling algorithm. We establish that any list scheduling algorithm has a performance guarantee of 2, whatever its priority rule. Our analysis relies on 2 immediate lower bounds, namely the conservation of the work and the critical path. Let us denote

$$P = \max_{i=1,\dots,m} \left\{ \sum_{j=1}^n p_{ij} \right\} \quad \text{and} \quad Q = \max_{j=1,\dots,n} \left\{ \sum_{i=1}^m p_{ij} + q_j \right\}$$

Clearly $L_{\max}^* \geq P$ and $L_{\max}^* \geq Q$. We have the following result :

Proposition 1 *Any list scheduling algorithm is a 2-approximation algorithm for problem $Om||L_{\max}$. More precisely, for any list schedule π , $L_{\max}(\pi) \leq P + Q$*

Proof. Consider a list schedule π , and let u be a job such that $L_u = L_{\max}(\pi)$. Without loss of generality, we can assume that the last operation of u is scheduled on the first machine. We consider 2 cases : either an idle-time occurs on \mathcal{M}_1 before the completion of job u , or not. If there is no idle time on \mathcal{M}_1 , then $L_u \leq P + q_u \leq P + Q$. Otherwise, let us denote by I the total idle time occuring on \mathcal{M}_1 before the completion time of job u . We have $L_u \leq P + I + q_u$. Notice that job u could not have been available

on machine \mathcal{M}_1 at any idle instant, otherwise, due to the principle of list scheduling algorithms, it would have been scheduled. As a consequence, an operation of job u is performed on another machine at every idle instant of \mathcal{M}_1 before C_u . Hence, we can bound the idle time I by the total processing time of job u . We have :

$$L_u \leq P + I + q_u \leq P + \sum_{i=1}^m p_{iu} + q_u \leq P + Q$$

We can conclude that in any case $L_{\max}(\pi) \leq P + Q \leq 2L_{\max}^*$ ■

Notice that good *a posteriori* performances can be achieved by a list scheduling algorithm, for instance if the workload P is large compared with the critical path Q . One natural question is whether some better approximation ratios can be obtained with particular lists. It is a folklore that minimizing the maximum lateness on one resource can be achieved by sequencing the tasks in non-increasing order of their delivery times. This sequence is known as Jackson's order. One can wonder if a list scheduling algorithm using Jackson's order as its list performed better in the worst case. The answer is negative. The following proposition states that the analysis of Proposition 1 is tight whatever the list.


Proposition 2 *No list scheduling algorithm can have a performance ratio less than 2 for problem $O2||L_{\max}$.*

Proof. Consider the following instance: we have 3 jobs to schedule on 2 machines. Jobs 1 and 2 have only one (non null) operation to perform, respectively on machine \mathcal{M}_1 and \mathcal{M}_2 . The duration of the operation is equal to a time units, where $a \geq 1$ is a parameter of the instance. Both delivery times are null. Job 3 has one unit operation to perform on each machine, and its delivery time is $q_3 = a$.

An optimal schedule sequences first Job 3 on both machines, creating an idle time at the first time slot, and then performs Jobs 1 and 2. That is, the optimal sequence is $(3, 1)$ on \mathcal{M}_1 and $(3, 2)$ on \mathcal{M}_2 . The maximum lateness is equal to $L_{\max}^* = a + 2$. Notice that this schedule cannot be obtained by a list scheduling algorithm, since an idle time occurs at the first instant while a job (either 1 or 2) is available. Indeed, it is easy to see that, whatever the list, either Job 1 or Job 2 is scheduled at time 0 by a list scheduling algorithm. As a consequence, Job 3 cannot complete before time $a + 1$ in a list schedule π . Hence, $L_{\max}(\pi) \geq 2a + 1$. The ratio for this instance is $\frac{2a+1}{a+2}$, which tends to 2 when a tends to $+\infty$. ■

As a conclusion, Jackson's list does not perform better than any other list in the worst case. Nevertheless, we use it extensively in the PTAS that we present in the next section.

Optimal schedule:



Jackson's schedule:




Figure 1. Illustration of the worst case of Jackson's rule

3 PTAS

In this section, we present the first PTAS for problem $Om||L_{\max}$, that is, when the number of machines is fixed. Our algorithm considers three classes of jobs as introduced by Sevastianov and Woeginger [10] and used by several authors for a variety of makespan minimization in shops (see for instance the extension by Jansen *et al.* for the job shop [5]). Notice that our approximation algorithm does not require to solve any linear program.

3.1 Description of the Algorithm

Let ε be a fixed positive number. We describe how to design an algorithm, polynomial in the size of the inputs, with a performance ratio of $(1 + \varepsilon)$ for problem $Om||L_{\max}$. As a shorthand, let $\bar{\varepsilon} = \frac{\varepsilon}{2m(m+1)}$. Recall that $P = \max_{i=1}^m \{\sum_{j=1}^n p_{ij}\}$ is the maximal workload of a machine. For a given integer k , we introduce the following subsets of jobs \mathcal{B} , \mathcal{S} and \mathcal{T} :

$$\mathcal{B} = \left\{ j \in \mathcal{J} \mid \max_{i=1}^m p_{i,j} \geq \bar{\varepsilon}^k P \right\} \quad (2)$$

$$\mathcal{S} = \left\{ j \in \mathcal{J} \mid \bar{\varepsilon}^k P > \max_{i=1}^m p_{i,j} \geq \bar{\varepsilon}^{k+1} P \right\} \quad (3)$$

$$\mathcal{T} = \left\{ j \in \mathcal{J} \mid \bar{\varepsilon}^{k+1} P > \max_{i=1}^m p_{i,j} \right\} \quad (4)$$

By construction, for any integer k , sets \mathcal{B} , \mathcal{S} and \mathcal{T} define a partition of the jobs. For the ease of understanding, the jobs of \mathcal{B} will be often called the *big* jobs, the jobs of \mathcal{S} the *small* jobs, and the jobs of \mathcal{T} the *tiny* jobs. Notice that the duration of any operation of a small jobs is less than $\bar{\varepsilon}^k P$, and less than $\bar{\varepsilon}^{k+1} P$ for a tiny job. The choice of k relies on the following proposition, which comes from Sevastianov and Woeginger [10]:

Proposition 3 [10] *There exists an integer $k \leq \lceil \frac{m}{\varepsilon} \rceil$ such that*

$$p(\mathcal{S}) \leq \varepsilon P \quad (5)$$

where $p(\mathcal{S}) = \sum_{j \in \mathcal{S}} \sum_{i=1}^m p_{ij}$ is the total amount of work to perform for the jobs of \mathcal{S} . Moreover, for the big jobs, we have:

$$|\mathcal{B}| \leq \frac{m}{\bar{\varepsilon}^k} \quad (6)$$

Proof. Let us denote $z = \lceil m/\varepsilon \rceil$. Observe that for a given value k , the duration of the largest operation of any small job belongs to the interval $I_k = [\bar{\varepsilon}^{k+1} P, \bar{\varepsilon}^k P[$. Assume for the sake of contradiction that, for all values $k = 1, \dots, z$, the corresponding set \mathcal{S}_k does not verify Condition (5). As

a consequence, $p(\mathcal{S}_k) > \varepsilon P$ for each $k = 1, \dots, z$. Since these sets are disjoint, it results that the total processing time of the operations of the jobs whose the duration of its largest operation belongs to $\in [\bar{\varepsilon}^{z+1}P, P[$ is strictly greater than $z\varepsilon P$. However, this amount of work is bounded by the total work of the instance. We have :

$$z\varepsilon P < \sum_{i=1}^m \sum_{j=1}^n p_{ij} \leq mP$$

Thus $z < m/\varepsilon$, which contradicts our definition of z . It follows that at least one interval $\mathcal{I}_k = [\bar{\varepsilon}^{k+1}P, \bar{\varepsilon}^k P[$ with $1 \leq k \leq z$ is suitable to contain the values of the large operations of subset \mathcal{S} such that $p(\mathcal{S}) \leq \varepsilon P$.

To prove Inequality (6), we can observe that the total processing time of the operations of \mathcal{B} is bounded by mP . Thus, $|\mathcal{B}| \bar{\varepsilon}^k P \leq mP$ must hold and Inequality (6) follows. ■

Notice that, for a fixed value m of machines, only a constant number $\lceil m/\varepsilon \rceil$ of values must be considered for k . Hence, an integer k verifying the conditions of Proposition 3 can be found in linear time. Assume from now that k has been chosen according to Proposition 3. In order to present our approach, let us explain how the different sets \mathcal{S} , \mathcal{B} and \mathcal{T} of jobs are scheduled in our PTAS. Since set \mathcal{S} represents a very small work, we can schedule it first. Clearly, its last operation cannot complete after time $t(\mathcal{S}) \leq \varepsilon P$ in a list schedule. Since set \mathcal{B} has a fixed number of jobs, we can afford to consider all the ways to sequence them. For that, we discretize the time, considering a time step $\delta = \bar{\varepsilon}^{k+1}P$. Finally, for each assignment of the big jobs, we schedule the tiny jobs using simply Jackson's list scheduling algorithm. One originality of our approach is the possibility for a tiny job to push a big job in order to fit before it. More precisely, if the tiny job the list scheduling algorithm is considering cannot complete before the start of the next big job on its machine, say b , then we force its schedule by shifting right the operation of job b as much as necessary. This shifting is special in twofolds : first, we also shift right of the same amount of time *all* the operations of the big jobs starting after job b . Second, the operation of job b is then *frozen*, that is, it cannot be pushed again by a tiny job. Hence, an operation of a big job can be pushed at most once by a tiny job, but can be shifted right a lot of times, due to the push of other operations of some big jobs. A more formal description of our algorithm can be given as follows:

ALGORITHM PTAS

1. Schedule first jobs of \mathcal{S} using any list scheduling algorithm between time 0 to time $p(\mathcal{S})$ (the cost factor of this simplification will not be more than $1 + \varepsilon$).
2. Let $\delta = \bar{\varepsilon}^{k+1}P$. Consider all the time intervals between $p(\mathcal{S})$ and mP of length δ (the number of these intervals is a constant for a fixed ε).

3. Enumerate all the schedules of jobs in \mathcal{B} between $p(\mathcal{S})$ and mP . Here, a schedule is reduced to an assignment of the operations to starting times of the time intervals defined in the previous step (the cost factor of this simplification will not be more than $1 + \varepsilon$).
4. Complete every partial schedule generated in the last step by adding the jobs of \mathcal{T} . The operations of \mathcal{T} are added by applying a list scheduling algorithm using Jackson's order (i.e., when several operations are available to be performed we start by the one of the largest delivery time). Note that if an operation cannot fit in front of a big job b , then we translate b and all the next big jobs by the same necessary duration to make the schedule feasible. The operation of job b is then frozen, and cannot be shifted any more.
5. Return the best feasible schedule found by the algorithm.

3.2 Analysis of the Algorithm

We start by introducing some useful notations. Consider a schedule π . For each machine i , we denote respectively by s_{ir} and e_{ir} the start time and completion time of the r th operation of a big job on machine i , for $r = 1, \dots, |\mathcal{B}|$. By convenience we introduce $e_{i0} = 0$. For short we call the *grid* the set of all the couples (resource \times starting time) defined in Phase (2) of the algorithm. Recall that in the grid the starting times are discretized to the multiples of δ . Notice that our algorithm enumerates in Phase (3) all the assignments of big job operations to the grid. Phase (4) consists in scheduling all the tiny jobs in-between the big jobs. In the following, we call a time-interval on a machine corresponding to the processing of a big job a *hole*, for the machine is not available to perform the tiny jobs. The duration of the r th hole on machine i , that is $e_{ir} - s_{ir}$, is denoted by h_{ir} . By analogy to packing, we call a *bin* the time-interval between two holes. The duration of the r th bin on machine i , that is $s_{ir} - e_{i,r-1}$, is denoted by a_{ir} . We also introduce $H_{ir} = h_{i1} + \dots + h_{ir}$ and $A_{ir} = a_{i1} + \dots + a_{ir}$, that is the overall duration of the r first holes and bins, respectively, on machine i .

Now consider an optimal schedule π^* . With immediate notations, let s_{ir}^* be the start time of the r th operations of a big job on the machine i , and let A_{ir}^* be the overall duration of the r first bins. For the ease of the presentation, we assume in the reminder, without loss of generality, that we have no small jobs to schedule : Indeed, Phase (1) does not increase the length of the schedule by more than $\varepsilon P \leq \varepsilon L_{\max}^*$. We say that an assignment to the grid is *feasible* if it defines a feasible schedule for the big jobs. The next lemma shows that there exists a feasible assignment such that each operation of the big jobs is delayed, compared to an optimal schedule, by at least $2m\delta$ time units and by at most $(2 + |\mathcal{B}|)m\delta$ time units.

Lemma 4 *There exists a feasible assignment \bar{s} to the grid such the operations of the big jobs are sequenced in the same order, and for every machine i and index r we have:*

$$s_{ir}^* + 2m\delta \leq \bar{s}_{ir} \leq s_{ir}^* + (2 + |\mathcal{B}|)m\delta$$

Proof. Among all the possible assignments enumerated in Phase (3) for the big jobs, certainly we consider the following one, which corresponds to a shift of the optimal schedule π^* restricted to the big jobs :

- Insert $2m\delta$ extra time units at the beginning of π^* , that is delay all the operations by $2m\delta$.
- Align the big jobs to the grid (shifting them to the right)
- Define the assignment \bar{s} as the current starting times of the operations of the big jobs.

More precisely, to align the big jobs to the grid, we consider sequentially the operations by non-decreasing order of their starting time. We then shift right the current operation to the next point of the grid and translate the subsequent operations of the same amount of time. This translation ensures that the schedule remains feasible for the big jobs.

By construction each operation is shifted right by at least $2m\delta$ time units, which implies that $\bar{s}_{ir} \geq s_{ir}^* + 2m\delta$. The alignment of an operation to the grid again shifts it right, together with all the subsequent operations, by at most δ time units. Thus, the last operation is not shifted more than $m|\mathcal{B}|\delta$ time units by the alignment. The result follows. ■

Now consider the schedule π obtained by applying the Jackson's list scheduling algorithm to pack the tiny jobs between the holes, starting from the feasible assignment of Lemma 4. Notice that, due to the shift procedure in Phase (4), the starting time of the big jobs (the holes) can change between the assignment \bar{s} and the schedule π . However a hole can be shifted at most $m|\mathcal{B}|$ times since each operation of a big job is shifted at most once by a tiny job. Moreover the length of a shift is bounded by the duration of an operation of a tiny job, that is by δ . In addition, as we shift all the operations belonging to the big jobs, the length of the bins cannot decrease in the schedule π . Hence, we have the two following properties for the schedule π , which are direct consequences of Lemma 4 and of the previous discussion :

1. Any operation of a big job is only slightly delayed compared to the optimal schedule π^* : $s_{ir} \leq s_{ir}^* + 2(|\mathcal{B}| + 1)m\delta$

2. Each bin is larger in π than in the optimal schedule. More precisely we have $A_{ir} \geq A_{ir}^* + 2m\delta$ for all machine i and all index r .

In other words in the schedule π we have slightly delayed the big jobs to give more room in each bin for the tiny jobs. We say that a job y is more critical than a job x if y has a higher priority in the Jackson's order. By convention a job is as critical at itself. We have the following lemma:

Lemma 5 *In schedule π , for every job x , there exists a job y such that :*

$$q_y \geq q_x \quad \text{and} \quad C_x \leq C_y^* + 2(|\mathcal{B}| + 1)m\delta$$

Proof. Let x be a job and C_x its completion time in schedule π . Without loss of generality we can assume that the last operation of the job x is processed on the first machine. If x is a big job, that is $x \in \mathcal{B}$, we have already noticed that we have $C_x \leq C_x^* + 2(|\mathcal{B}| + 1)m\delta$, due to our choice of the big jobs assignment on the grid. Hence, the inequality of Lemma 5 holds for x . Thus consider in the remaining of the proof the case of a tiny job x . We denote by \mathcal{T}_x the subset of tiny jobs that are more critical than x and such that their operation on the first machine is completed by time C_x , that is :

$$\mathcal{T}_x = \{y \in \mathcal{T} \mid C_y^1 \leq C_x^1 \text{ and } q_y \geq q_x \}$$

Observe that our definition implies in particular that $x \in \mathcal{T}_x$. We first establish that in schedule π , almost all the tiny jobs processed before x on the first machine are more critical than x . That is, the schedule π essentially follows the Jackson's sequence for the tiny jobs. Let r be the index of the bin where x completes in schedule π . For short we denote by $A_1(x)$ the overall time available for processing tiny jobs on the first machine over the time-interval $[0, C(x)]$, that is $A_1(x) = C_x - H_{1,r}$. We also denote by $p_1(\mathcal{T}_x)$ the total processing time of the operations of \mathcal{T}_x on the first machine. We claim that :

$$p_1(\mathcal{T}_x) \geq A_1(x) - 2(m-1)\delta \tag{7}$$

If at every available instant on the first machine till the completion of x an operation of \mathcal{T}_x is processed in π , then clearly we have $A_1(x) = p_1(\mathcal{T}_x)$ and Inequality 7 holds. Otherwise, consider a time interval $I = [t, t']$, included in a bin, such that no task of \mathcal{T}_x is processed. We call such an interval *non-critical* for x . It means that during I , either some idle times appear on the first machine, and/or some jobs less critical than x have been processed. However, due to the shift procedure and the Jackson's list used by the algorithm, the only reason for not scheduling x during I is that this job is not available by the time another less critical job z is started. Notice that in an open-shop environment, a job x is not available on the first machine only if one of its operations is being processed on another machine. As a

consequence, the interval I necessarily starts during the processing of x on another machine, that is $t \in [t_{i,x}, C_{i,x}]$ for some machine i . This holds for any idle instant and any time an operation is started in interval I . As a consequence, the interval I cannot finish later than the completion of x on another machine i' , plus the duration of a less critical (tiny) job z eventually started on the first machine during time interval $[t_{i',x}, C_{i',x}]$. Since all the jobs are tiny, the overall duration of the non-critical intervals for x is thus bounded by $\sum_{i=2}^m (p_{i,x} + \delta)$, which is at most equal to $2(m-1)\delta$. Inequality 7 follows.

Now let y be the job of \mathcal{T}_x that completes last on the first machine in the optimal schedule π^* . Let r^* be the index of the bin where y is processed on the first machine in π^* , and let $A_1^*(y)$ be the total available time for tiny jobs in π^* before time $C_{1,y}^*$, that is $A_1^*(y) = C_{1,y}^* - H_{1,r^*}^*$. Recall that r is the number of bins used in schedule π to process all the operations of \mathcal{T}_x on the first machine. We prove that the optimal schedule also uses (at least) this number of bins, that is $r^* \geq r$. Indeed, by the conservation of work, we have that $A_1^*(y) \geq p_1(\mathcal{T}_x)$. Using Inequality 7 we obtain that $A_1^*(y) \geq A_1(x) - 2(m-1)\delta$. By definition of r and r^* , we also have $A_{1,r-1} \leq A_1(x)$ and $A_1^*(y) \leq A_{1,r^*}^*$. Hence, the following inequality must hold:

$$A_{1,r-1} \leq A_{1,r^*}^* + 2(m-1)\delta$$

However, we have observed that our choice of the assignment of the big jobs to the grid ensures that for any index l , $A_{1,l}^* + 2m\delta \leq A_{1,l}$, which implies that we have $A_{1,r-1} + 2\delta \leq A_{1,r^*}$. As a consequence, inequality $A_{1,r-1} < A_{1,r^*}$ must hold. Since $A_{1,l}$ represents the total length of the l first bins in π , which is obviously non-decreasing with l , it implies that $r \leq r^*$.

It means that in π^* , task y cannot complete its operation on the first machine before the first r big tasks. We can conclude the proof of Lemma 5 by writing that, on one hand, $C_y^* \geq p_1(\mathcal{T}_x) + H_{1,r}$, and, on the other hand, $C_x \leq p_1(\mathcal{T}_x) + 2(m-1)\delta + H_{1,r}$. As a consequence, x does not complete in π later than $2(m-1)\delta$ times units after the completion time of y in π^* . Since by definition y is more critical than x , Lemma 5 follows. ■

Finally, we can conclude that the following theorem holds:

Theorem 6 *Problem $Om||L_{\max}$ admits a PTAS.*

Proof. We first establish that the maximum lateness of the schedule returned by our algorithm is bounded by $(1 + \varepsilon)L_{\max}^*$. In schedule π defined in Lemma 5, let u be a job such that $L_{\max}(\pi) = C_u + q_u$. If job u is a small job, then it completes before time $p(\mathcal{S})$. Due to our choice of the partition, see Proposition 3, we have $L_u \leq \varepsilon P + q_u \leq (1 + \varepsilon)L_{\max}^*$. Hence, in the following, we restrict to the case where $u \notin \mathcal{S}$, that is, job u is either a big

or a tiny job. According to Lemma 5, there exists a job y such that :

$$q_u \leq q_y \quad \text{and} \quad C_u \leq C_y^* + 2(|\mathcal{B}| + 1)m\delta$$

We have :

$$\begin{aligned} L_{\max}(\pi) &= C_u + q_u \\ &\leq C_y^* + 2(|\mathcal{B}| + 1)m\delta + q_y \\ &\leq L_{\max}^* + 2(|\mathcal{B}| + 1)m\delta \end{aligned}$$

As a consequence, using Proposition 3, we can write that for any fixed $\varepsilon \leq 1$:

$$\begin{aligned} L_{\max}(\pi) - L_{\max}^* &\leq 2\left(\frac{m}{\varepsilon^k} + 1\right)m\varepsilon^{k+1}P \\ &\leq 2\left(\frac{m+1}{\varepsilon^k}\right)m\varepsilon^{k+1}P \\ &= 2(m+1)m\varepsilon P \\ &\leq \varepsilon L_{\max}^* \end{aligned}$$

Hence, our algorithm has a performance guarantee of $(1 + \varepsilon)$. Let us now check its time complexity. First, the identification of k and the three subsets \mathcal{B} , \mathcal{S} and \mathcal{T} can be done in $O(\frac{m^2 n}{\varepsilon})$. Second, the scheduling of the jobs of \mathcal{S} can clearly be performed in polynomial time (in fact, in linear time in n for m fixed). Now, let us consider the scheduling of the big jobs. The number Δ of points in the grid is bounded by:

$$\begin{aligned} \Delta &\leq m \times \frac{mP}{\delta} = \frac{m^2}{\varepsilon^{k+1}} \\ &\leq \frac{m^2}{\varepsilon^{2+\frac{m}{\varepsilon}}} \\ &\leq m^2 \left(\frac{2m(m+1)}{\varepsilon} \right)^{2+\frac{m}{\varepsilon}} \end{aligned}$$

The second inequality comes from the fact that $k \leq \lceil m/\varepsilon \rceil$, due to Proposition 3. The last bound is clearly a constant for m and ε fixed. The number of possible assignments of jobs of \mathcal{B} in Phase (3) is bounded by

$$(m|\mathcal{B}|)^\Delta \leq \left(\frac{m^2}{\varepsilon^k} \right)^\Delta$$

Hence, only a constant number of assignments to the grid are to be considered. Phase (4) completes every feasible assignment in a polynomial time. Phase (5) outputs the best solution in a linear time of the number of feasible assignments. In overall, the algorithm is polynomial in the size of the instance for a fixed ε and a fixed m . ■

4 Conclusion

In this paper we considered an open question related to the existence of PTAS to the m -machine open shop problem where m is fixed and the jobs have different delivery times. We answered successfully to this important question. This represents the best possible result we can expect due to the strong NP-hardness of the studied problem.

Our perspectives will be focused on the study of other extensions. Especially, the problem with release dates seems to be very challenging.

References

- [1] Averbakh I., Berman O., Chernykh I., 2005. A $\frac{6}{5}$ -approximation algorithm for the two-machine routing open shop problem on a 2-node network. *European Journal of Operational Research* 166(1): 3-24.
- [2] Chernykh I., Kononov A., Sevastianov S., 2012. Efficient approximation algorithms for the routing open shop problem. *Computers & Operations Research*
- [3] Hall L.A., Shmoys D.B., 1992. Jackson's rule for single machine scheduling: making a good heuristic better. *Mathematics of Operations Research* 17: 22:35.
- [4] Gonzales T., Sahni S., 1976. Open shop scheduling to minimize finish time. *Journal of Association Computing Machinery* 23: 665-679.
- [5] Jansen K., Solis-Oba R., Sviridenko M., 2003. Makespan minimization in job-shops: A linear time approximation scheme. *SIAM Journal on Discrete Mathematics* 16(2): 288-300.
- [6] Kellerer H (2004) Minimizing the maximum lateness. In: *Leung JYT (ed) Handbook of scheduling: algorithms, models and performance analysis*, CRC press, chap. 10.
- [7] Kononov A., Sviridenko M., 2002. A linear time approximation scheme for makespan minimization in an open shop with release dates. *Operations Research Letters* 30: 276-280.
- [8] Lawler E.L., Lenstra J.K., Rinnooy Kan A.H.G., 1981. Minimizing Maximum Lateness in a Two-Machine Open Shop. *Mathematics of Operations Research* 6:153-158.
- [9] Lawler E.L., Lenstra J.K., Rinnooy Kan A.H.G., 1982. Erratum. *Mathematics of Operations Research* 7(4):635-635. <http://dx.doi.org/10.1287/moor.7.4.635>

- [10] Sevastianov S.V., Woeginger G.J., 1998. Makespan minimization in open shops: A polynomial time approximation scheme. *Mathematical Programming* 82: 191-198.
- [11] Williamson D.P. , Hall L.A., Hoogeveen J.A., Hurkens C.A.J., Lenstra J.K., Sevastianov S.V., Shmoys D.B., 1997. Short shop schedules. *Operations Research*, 45:288-294